

**NANYANG
TECHNOLOGICAL
UNIVERSITY**
SINGAPORE

SC3020 Database System Principles

AY23/24 Semester 1

Project 2 Report

Group 48

Team Members	
Atharv Gupta	U2123893E
Esther	U2122073H
Sim Jia Ming	N2304339G
Stephen	U2140974G

SC3020 Database System Principles	1
Acknowledgment	3
Individual Contributions	3
Objectives	4
Implementation and Results	5
Key Functions and Algorithms	12
draw_nodes_recursively	12
get_block_content	12
get_block_accessed_content	12
detect_injection	13
Limitations	14
Installation guide	14
Source Code	15
Software Requirements	15
Installation Steps	15

Acknowledgment

This project incorporates the use of a Large Language Model, specifically ChatGPT developed by OpenAI. We acknowledge and appreciate the contribution of OpenAI in providing access to this advanced language model, which has played a crucial role in assisting with the development of our user interface.

Individual Contributions

Name	Contributions
Atharv Gupta	1) Report
Esther	1) Implemented function to get user input for SQL query, table selection and blockID 2) Implemented visualisation of disk block content and access functions 3) Implemented UI frames for combined visualisation of part a) and b)
Sim Jia Ming	1) Implemented explore.py, project.py, database.py 2) Assisted with interface.py bugs fixes / implementation 3) Report - Installation guide, Architecture diagram, Key Functions and Algorithms
Stephen	1) Draw QEP Tree Functions 2) Display Node Details Functions 3) Integrated UI Functions

Objectives

There are two primary objectives of this project, initially it seeks a SQL query input from the user, upon receiving the input, the program creates a representation of the QEP in the form of an interactive operator tree, while also giving visual information about the disk blocks accessed while processing the query. Secondly it offers a user-friendly interface, which is the GUI component of the project, wherein the user can explore the contents of the disk blocks, by inputting the relation name and the block ID.

For this program, the host language has been chosen to be Python, while the Database Management System used is PostgreSQL. PostgreSQL is an open source relational database management system which in this project stores the TPC-H dataset, utilised for testing of the program.

The following sections of this report will give a basic idea of how this program has been implemented and the results displayed by it.

Implementation and Results

The implementation (as shown in figure 1) of this project has three major components which are:

- Interface: The GUI component, presenting a user friendly interface to input queries and observe the results
- Explore: The part of the code that explores the PostgreSQL database and fetches the information.
- Project: This is the executable file of the program.

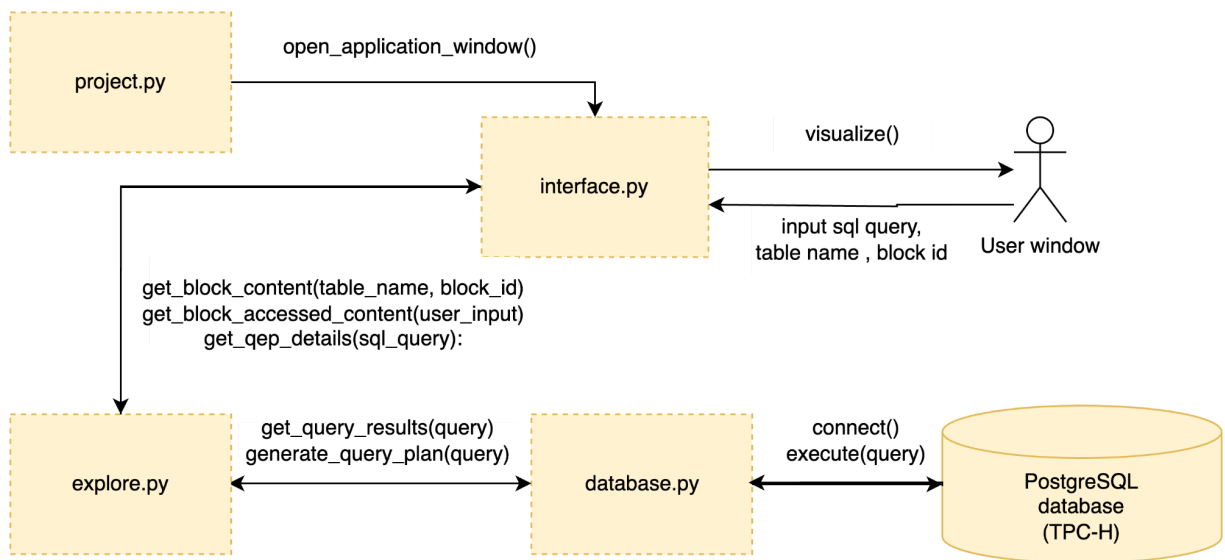


Figure 1. Architecture Diagram

User-Interface

Figure 2 illustrates the overview of the display a user would see after running the project.py file. The window would be titled “Database Block Visualisation” and will have five sections: Enter SQL Query, Block Accessed Information Table, Block Content, displaying the QEP and displaying the Node Details.

The screenshot shows a window titled "Database Block Visualization" with standard window controls (minimize, maximize, close). The interface is divided into five main sections:

- Enter SQL Query:** A text input field with a "Submit" button to its right.
- Block Accessed Information Table:** A large, empty rectangular box intended for a table.
- Block Content:** A section containing a dropdown menu labeled "Select a table:" with "customer" selected, an "Enter Block ID:" text input field, and a "Visualize" button.
- QEP:** A large, empty rectangular box on the left side of the bottom section.
- Node Details:** A large, empty rectangular box on the right side of the bottom section, featuring a vertical scrollbar on its right edge.

Figure 2: Overview of the Interface

- **SQL Query Input**

The user is expected to populate the “Enter SQL Query” box by inputting an SQL query of their choice, as shown in figure 3. Upon entry and reviewing, the user can click on the Submit button to move forward.

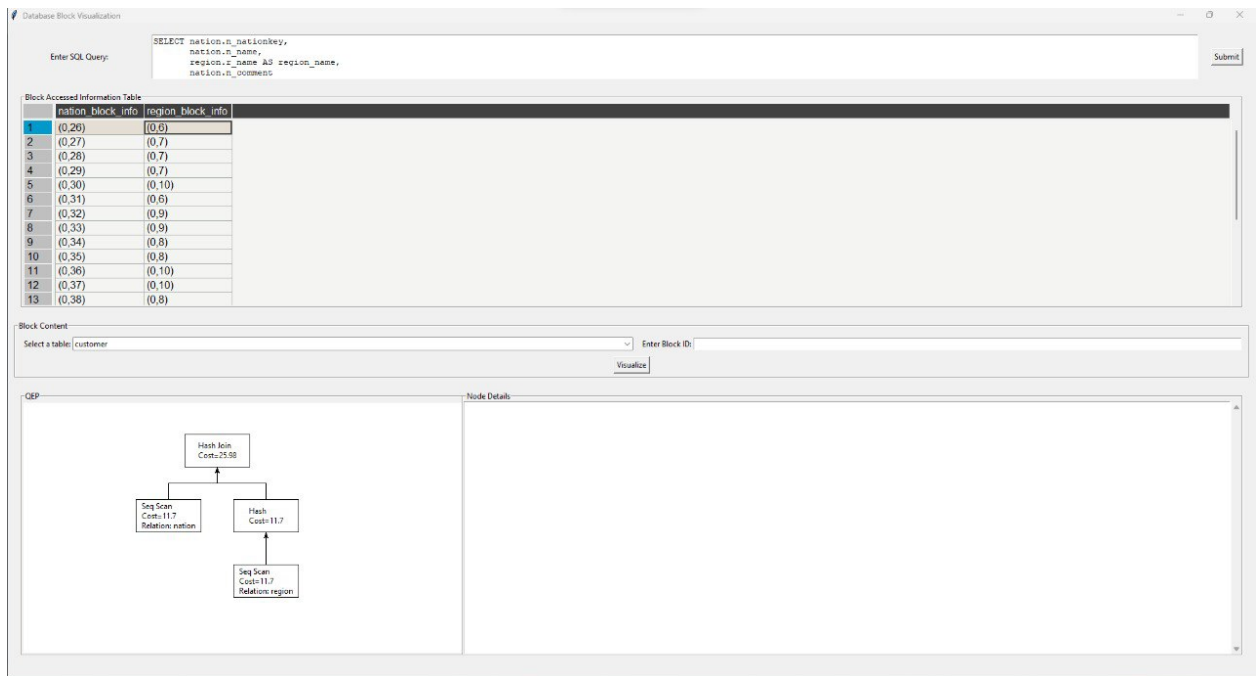


Figure 3: Entering SQL Query

In case, the query inputted is not valid or accurate, an error prompt as seen in figure 4 will be displayed. The prompt will show the error that the user had committed and upon dismissing the prompt can re-enter the corrected SQL query.




Figure 4: Error Message Due to Invalid Query

- **Visualisation of the Operator Tree**

Upon entry of a valid SQL query, the explore.py component of the program will process the query. Upon processing it will generate an interactive operator tree as seen in figure 5.

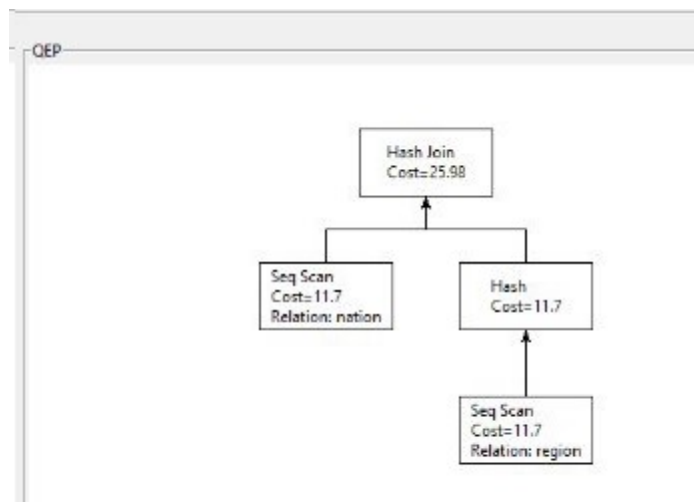


Figure 5: The QEP Operator Tree

The user can observe various operators, their ordering and execution during the query execution. The users can also click any node of the operator tree to get more information, as seen in figure 6.

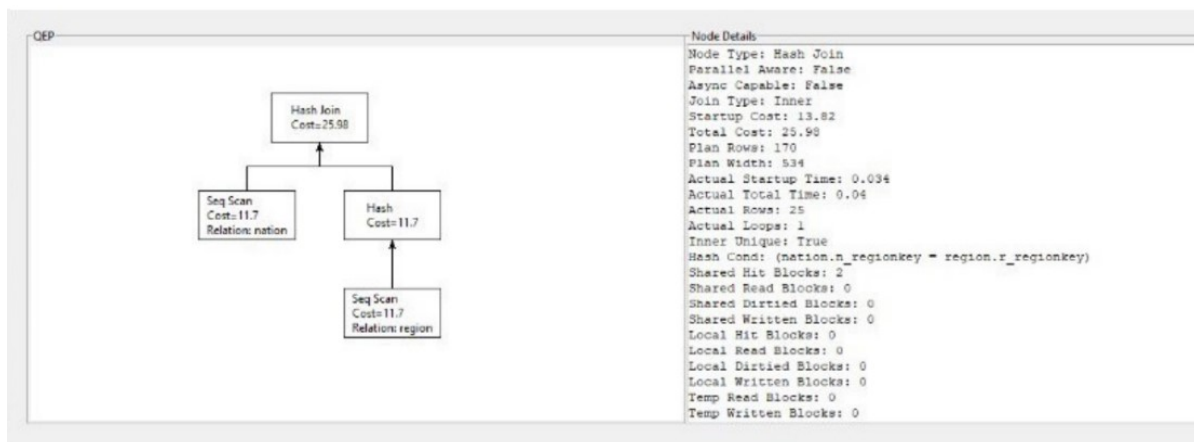


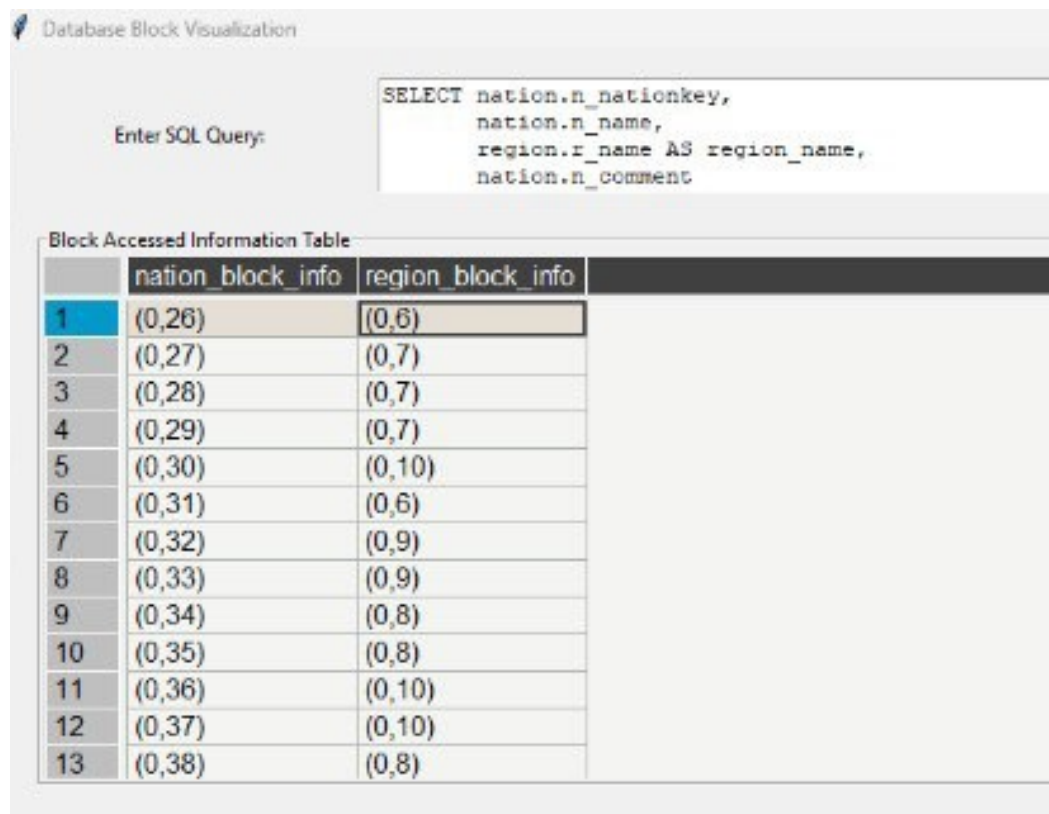
Figure 6: Node Details of Hash Join

In the “Node Detail” section, the user can gain further information about the operators used in every step, while also getting information about the number of

disk blocks accessed, total cost and other relevant information such as buffer hits. This is the representation of QEP and lets the user see the internal working of how a database executes a query. Upon changing the initially inputted query the operator tree will be displaced by the new desired tree.

- **Block Accessed Information Table**

The user also will be able to see the (block number, offset) pairs accessed by the query results in the “Block Accessed Information Table” section as seen in figure 7. These will be tabulated according to the table they belong to and are shown in this section.



The screenshot shows a web interface titled "Database Block Visualization". It has a text input field labeled "Enter SQL Query:" containing the following SQL query:

```
SELECT nation.n_nationkey,
       nation.n_name,
       region.r_name AS region_name,
       nation.n_comment
```

Below the query input is a section titled "Block Accessed Information Table" which contains a table with 13 rows and 3 columns. The first column is an index from 1 to 13. The second column is labeled "nation_block_info" and the third column is labeled "region_block_info". Each row contains a pair of (block number, offset) for the nation table and the region table respectively.

	nation_block_info	region_block_info
1	(0,26)	(0,6)
2	(0,27)	(0,7)
3	(0,28)	(0,7)
4	(0,29)	(0,7)
5	(0,30)	(0,10)
6	(0,31)	(0,6)
7	(0,32)	(0,9)
8	(0,33)	(0,9)
9	(0,34)	(0,8)
10	(0,35)	(0,8)
11	(0,36)	(0,10)
12	(0,37)	(0,10)
13	(0,38)	(0,8)

Figure 7: Block Accessed Information Table

So for example, the first tuple which shows (0,26) in the nation table and (0,6) in the region table are the block and tuple pair of the results of the query inputted by the user.

- **Visualisation of Disk Blocks**

The “Block Content” section of the interface, would allow the user to select a table name from a drop down list and enter the block ID to view the contents of that particular block. Upon entry and review, the user can click the “Visualise” button to view all the tuples situated in that block in the form of a pop up window as seen in figure 8.

c_custkey	c_name	c_address	c_nationkey	c_phone
43	Customer#00000043	ouSbjHk8lh5fKX3zGso3Z5j9Aa3PoaFd	19	29-316-665-2897
44	Customer#00000044	OiNdOSPwDu4jo4xNNP85E0dmhZGvNtBwi	16	26-190-260-5375
45	Customer#00000045	4v3OcpFgoOmMGNcBnFN4mdC	9	19-715-298-9917
46	Customer#00000046	eaTXWWm10L9	6	16-357-681-2007
47	Customer#00000047	b0UgocSqEW5 gdVbhNT	2	12-427-271-9466
48	Customer#00000048	0UU iPhBupFvemNB	0	10-508-348-5882
49	Customer#00000049	cNgAeX7FqrdF7HQN9EwjUa4nxTN68L FKAxzl	10	20-908-631-4424
50	Customer#00000050	9SzDYlkzxByyJ1QeTl o	6	16-658-112-3221
51	Customer#00000051	uRNwEaiTvo4	12	22-344-885-4251
52	Customer#00000052	7 QOqGqq5y9jfv51BC71jcHJSD0	11	21-186-284-5998
53	Customer#00000053	HnaxHzTffTZs8MuCplyTbZ47Cm4wFOOgib	15	25-168-852-5363
54	Customer#00000054	Nk4vf 5vECGWfYnHosTEN	4	14-776-370-4745
55	Customer#00000055	zIRBR4KNEI HzaiV3a i9n6elrxzDEh8r8pDom	10	20-180-440-8525
56	Customer#00000056	BJYZYJQk4yD58	10	20-895-685-6920
57	Customer#00000057	97XYbsuOPRXPWU	21	31-835-306-1650
58	Customer#00000058	g9ap7Dk1Sv9fcXEWjpMYpBZIRUohi T	13	23-244-493-2508
59	Customer#00000059	zLOCp0wh92OtBihgspOGI4	1	11-355-584-3112
60	Customer#00000060	FyodhjwMChsZmUz7Jz0H	12	22-480-575-5866
61	Customer#00000061	9kndve4EAJxhg3veF BfXr7AqOsT39o gtqjaYE	17	27-626-559-8599
62	Customer#00000062	upJK2Dnw13N	7	17-361-978-7059
63	Customer#00000063	IXRSpVWWZrakII	21	31-952-552-9584
64	Customer#00000064	MbCeGY20kaKK3oaUDNOT	3	13-558-731-7204
65	Customer#00000065	RGT yzQ0y4l0H90P783LG4U95bXQFDRXbWa1sINX	23	33-733-623-5267
66	Customer#00000066	XbsEqXH1ETbJYYtA1A	22	32-213-373-5094
67	Customer#00000067	rfG0cOgtr5W8 xILkwp9fpCS8	9	19-403-114-4356
68	Customer#00000068	o8AibcCRkXvQFh8hFN7o	12	22-918-832-2411
69	Customer#00000069	Ltx17nO9Wwhtdbe9QZVxNgP98V7xW97uvSH1prEw	9	19-225-978-5670
70	Customer#00000070	mFowluhnHjp2GjCiYYavkW kUwOjlaTCQ	22	32-828-107-2832
71	Customer#00000071	TlGalgdXWBmMVN6agLyWYDylz9MKzcY8glNw6t1B	7	17-710-812-5403
72	Customer#00000072	putjlmksxENzsNHqelA9Wqu7dhgH5BVCwDwHHcf	2	12-759-144-9689
73	Customer#00000073	8lhlxreu4Ug6tt5mog4	0	10-473-439-3214
74	Customer#00000074	lkJHCA3ZThF7qL7VKcrU nRLINKyif	4	14-199-862-7209
75	Customer#00000075	Dh 6jZNcwxWLKQfRKkiGrzv6pm	18	28-247-803-9025
76	Customer#00000076	m3sbCvjM0HyaOofHNe UkGPtqc4	0	10-349-718-3044
77	Customer#00000077	4tAE5KdMFGD4byHtXF92vx	17	27-269-357-4674
78	Customer#00000078	HB0taNZNqpg3U2cSL0kbrftkPwzX	9	19-960-700-9191

Figure 8: Visualisation of Disk Block 1 of customer table

If a user enters an invalid Block ID an error would be shown as in figure 9. Upon changing the inputted values, the initial set of tuples would be replaced by the desired table.

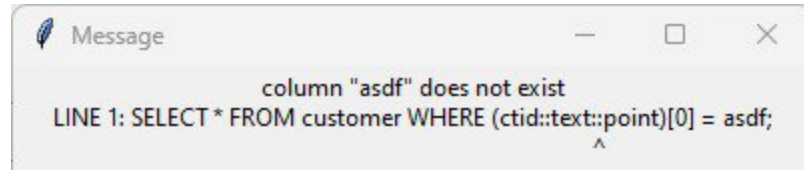


Figure 9: Error Message for Invalid Block ID

Accessory Files

- The project.py is the executable file of the program. It executes the software, by invoking necessary functions, handling input/output of the queries and visualisation information.
- The database.py on the other hand, ensures connectivity of the program with the PostgreSQL DBMS.
- The test.py is a file made for testing this program.

Key Functions and Algorithms

draw_nodes_recursively

This function is used to visually represent a QEP in a tree structure on the canvas. It extracts key node information like type, cost, and relation name from the QEP in JSON format. Each node is drawn as a rectangle on the canvas and clicking on the node would trigger an event to display the detailed information of the node, thus allowing the user to explore and visualise different aspects of the QEP. If a node has a parent it would connect the child nodes to their parents with a line. Using a depth-first search approach, the function would be recursively called if the current node has child nodes, allowing for an interactive visualisation of the QEP.

get_qep_details

This function returns the QEP in JSON format. Given an SQL query user input for example `SELECT * FROM region`, we prepend `EXPLAIN (ANALYZE, BUFFERS, FORMAT JSON)` to the SQL user input so that we can get the query execution plan from PostgreSQL in JSON format.

get_block_content

Given the table name and block id from user input, this function returns a table containing all the tuples in that table and block requested by the user. Using the table name and block id, we craft the following query `SELECT * FROM {table name} WHERE (ctid::text::point)[0] = {block id}` to retrieve the block content.

get_block_accessed_content

This function returns the (block number, offset) pairs accessed by the query results as explained in the “Block Accessed Information Table” section. Given an SQL query input from the user, we first process the input and check if there is any SQL injection.

Next, for the ease of parsing later, we convert the query to lowercase, remove all trailing spaces, semicolons, excess white spaces, new lines and tabs.

Subsequently, we check if this process input contains nested queries as this function does not support inputs with nested queries (limitation).

Following up, we removed all the `GROUP BY`, `HAVING` and all the aggregation functions (`SUM`, `AVY`, `MIN`, `MAX`, `COUNT`) as they affect the results shown in the Block Accessed Information Table.

Afterwards, we scanned the query to look for all the relations involved and all their respective `ctid` columns into the query.

Lastly, we crafted the query which is then executed to get the table of blocks accessed by the results of the user SQL query input.

Example user input:

```
SELECT * FROM region;
```

Example crafted query:

```
SELECT      region_block_info      FROM      (SELECT      region.ctid      as  
region_block_info, * FROM region) AS results_table;
```

detect_injection

This function is designed to detect and prevent CUD (Create, Update, Delete) types of SQL injection attempts. SQL injection is a technique where an attacker inserts malicious SQL code into a query, potentially leading to unauthorised access or manipulation of the database. We have blocked the following operations: `DELETE`, `UPDATE`, `INSERT`, `CREATE`, `DROP`, `ALTER`, `INDEX`.

Limitations

Our program has the following limitations:

1. **Disk-block access:** Our program cannot visualise all the disk blocks that have been accessed by an input query as `get_block_accessed_content` function only retrieves the disk blocks that contain the result of the query. This is a subset of all disk blocks that were accessed in order to get the result. However, one can view the block information of the results retrieved by the user query input with the help of the Block Accessed Information Table. Moreover, the user can view the number of disk blocks accessed upon clicking the root node of the operator tree.
2. **Block Accessed Information Table for Nested Queries:** Our program cannot display the information of the blocks accessed when the user enters a nested query. It would show an error message as shown in Figure 10

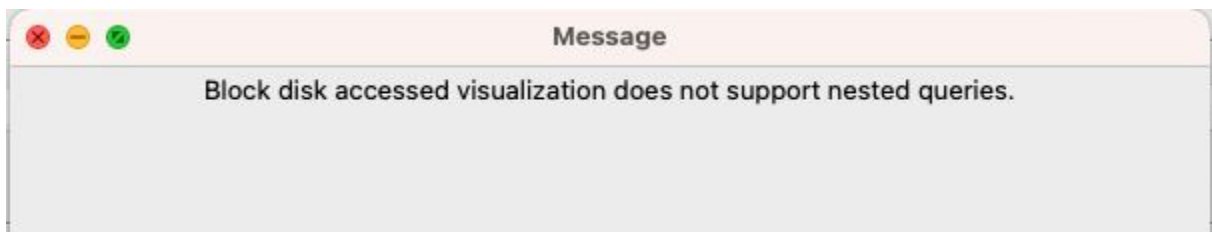


Figure 10: Error Message for Nested Queries

Installation guide

Source Code

- GitHub Repository - <https://github.com/SimJM/Database-System-Project-2>

Software Requirements

Before you start, make sure you have the following software installed:

- PyCharm - <https://www.jetbrains.com/pycharm/download/>
- Python 3.11 - <https://www.python.org/downloads/>
- PostgreSQL - <https://www.postgresql.org/download/>
- Git - <https://github.com/git-guides/install-git>

Installation Steps

Follow these steps to clone, configure, and run the project on your system:

Step 0: Set up PostgreSQL and create TPC-H

Follow the steps in the appendix of the project document.

Step 1: Clone the repository

Open your command prompt and run the following command to clone the project repository. Alternatively, you can simply download the zip file.


```
git clone https://github.com/SimJM/Database-System-Project-2.git
```

Step 2: Open the project in PyCharm

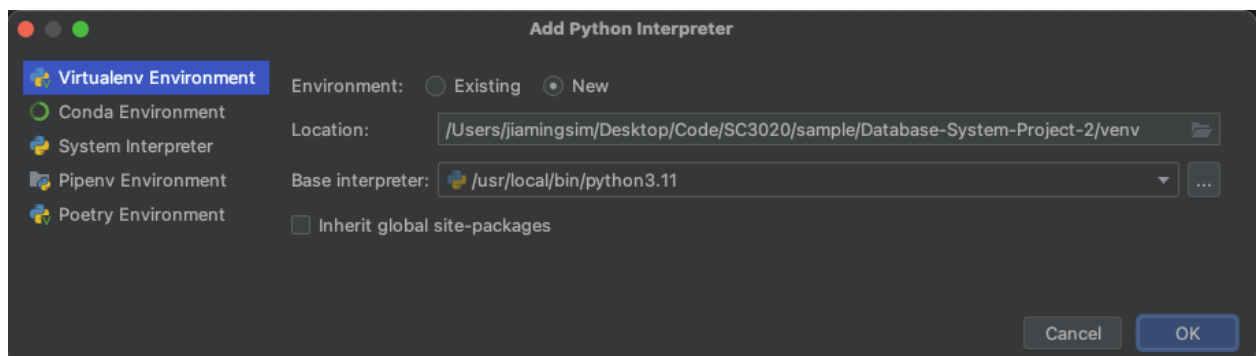
Launch PyCharm and open the `Database-System-Project-2` directory.

Step 3: Configure Python Interpreter

Open `project.py` then look for Configure Python interpreter to configure as shown below.



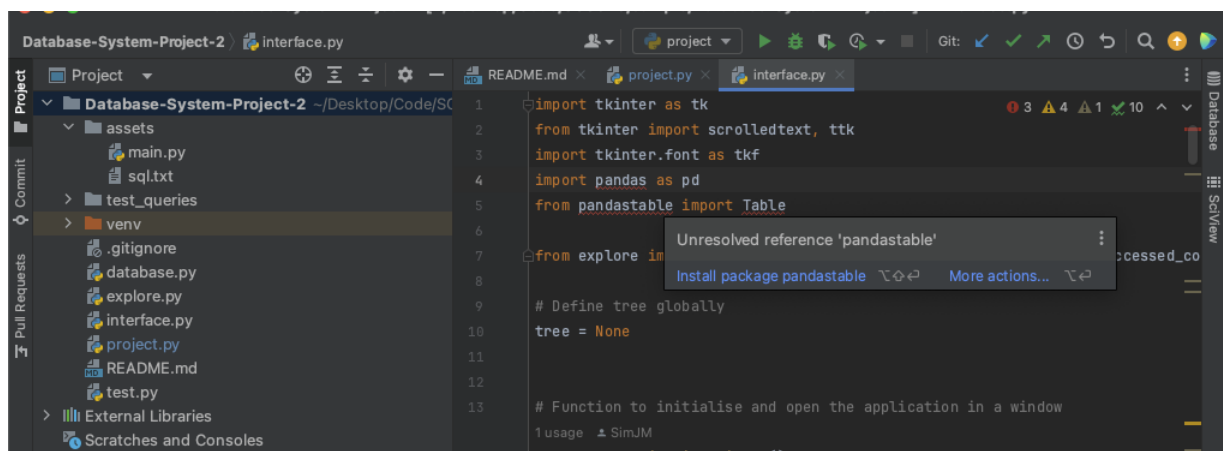
Example configuration:



Step 4: Install the relevant packages.

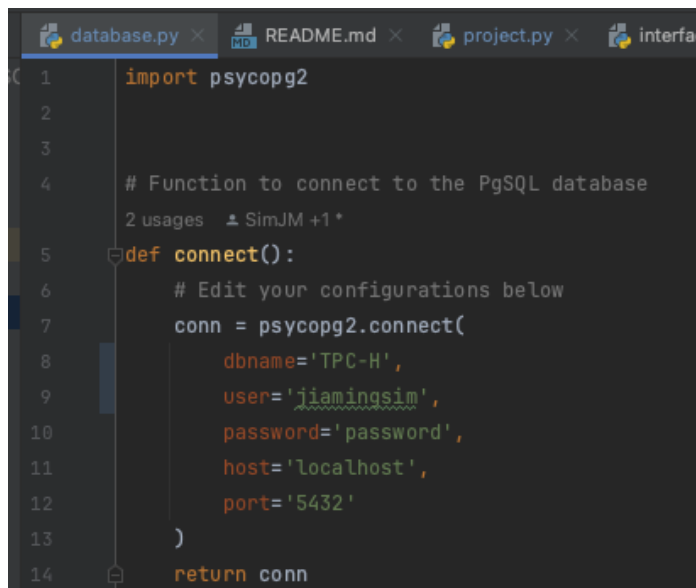
Open `project.py`, `explore.py`, `interface.py`, `database.py` and look for all the underlined imports in red. Right click on the underlined word then install package `<x>` as shown below.

Example:



Step 5: Edit PostgreSQL database credentials

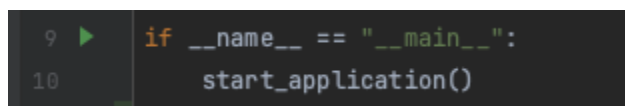
Open `database.py` and edit your PostgreSQL credentials.



```
1 import psycopg2
2
3
4 # Function to connect to the PostgreSQL database
5 2 usages  + SimJM +1 *
6 def connect():
7     # Edit your configurations below
8     conn = psycopg2.connect(
9         dbname='TPC-H',
10        user='jiamingsim',
11        password='password',
12        host='localhost',
13        port='5432'
14    )
15 return conn
```

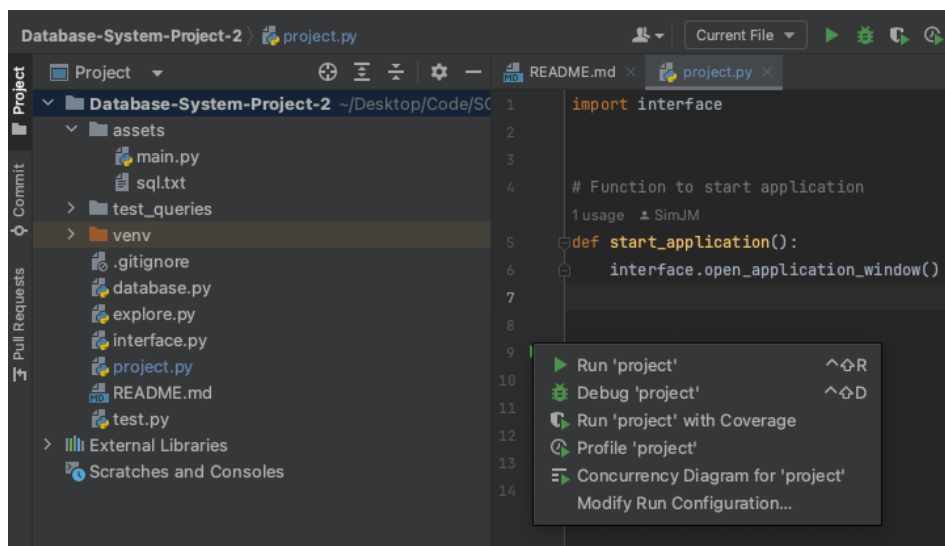
Step 6: Run the project

Once all the packages are installed, open `project.py` and look for the 'play' button in line 9.



```
9  ▶ if __name__ == "__main__":
10      start_application()
```

Click Run 'project' as shown below.



The screenshot shows the IDE interface with the `project.py` file open. The code in the editor is:

```
1 import interface
2
3
4 # Function to start application
5 1 usage  + SimJM
6 def start_application():
7     interface.open_application_window()
8
9
10
11
12
13
14
```

A context menu is open over the code, showing the following options:

- Run 'project' ^⇧R
- Debug 'project' ^⇧D
- Run 'project' with Coverage
- Profile 'project'
- Concurrency Diagram for 'project'
- Modify Run Configuration...

You should see the application open up in a window.

The screenshot shows a web application window titled "Database Block Visualization". The interface includes a text input field for "Enter SQL Query:" with a "Submit" button to its right. Below this is a large, empty rectangular box labeled "Block Accessed Information Table". Underneath that is a section titled "Block Content" which contains a dropdown menu labeled "Select a table:" with "customer" selected, and an "Enter Block ID:" text input field. A "Visualize" button is positioned below these two inputs. The bottom of the window is divided into two side-by-side panels: "QEP" on the left and "Node Details" on the right, both of which are currently empty.